


DD2421

Introduction to Artificial Neuronal Networks


Jörg Conradt

NeuroComputing Systems
Computational Science and Technology
School of Electrical Engineering and Computer Science
KTH Stockholm, Sweden

<http://neurocomputing.systems>



NCS




[HTTP://NEUROCOMPUTING.SYSTEMS](http://neurocomputing.systems)
CST - EECS - KTH STOCKHOLM, SWEDEN

Overview

- Biological Neurons
- Artificial Neurons
- What can a single Neuron do?
- From Neurons to Networks
- How to Train a Neuronal Network
- Tricks of the Trade ... how to get your Network working
- Software Example
- Other classes of Neuronal Network

Biological Neurons





electrically excitable cells for
processing and transmitting
information
(ca. 100 billion in the human brain)

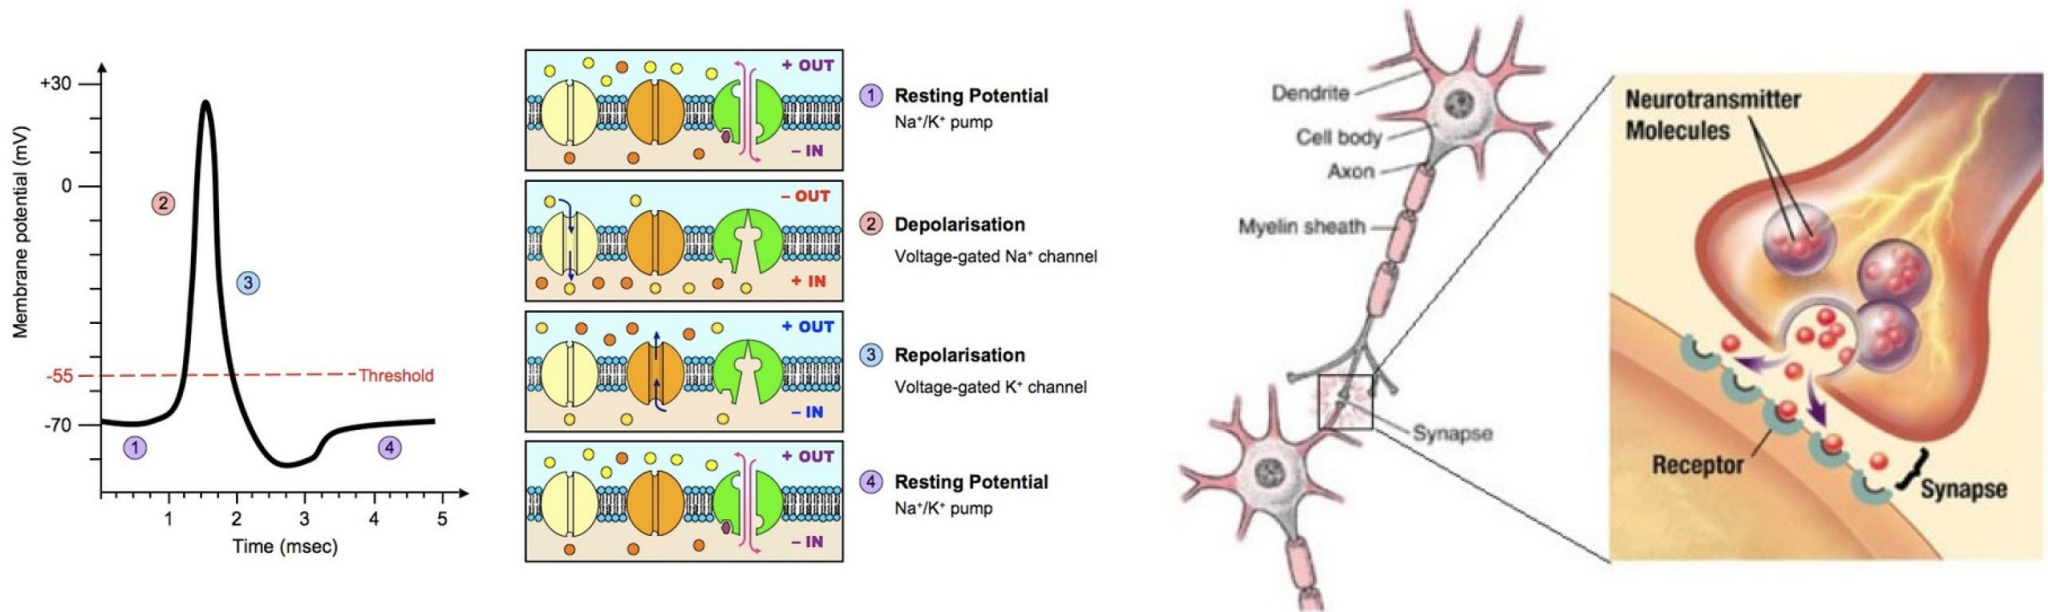
Biological Neurons

A large variety of cell types, but similar function:

(leaky) Integrate and Fire



Biological Details






More details on [neuroscience](https://www.kth.se/social/course/DD2401/)? KTH CST course in P4

<https://www.kth.se/social/course/DD2401/>

Artificial Neuron


From Biology to Technology




- Analog Leaky Integration
- Asynchronous Operation
- Spiking Communication
- Energy Efficient Implementation

?

Artificial Neuron



- Discrete Summation of Weighted Input
- Centrally Clocked System
- Continuous Values Communication (events \rightarrow rates)
- Von Neumann Computing Architectures




We transition to a technical abstraction, the *Perceptron*, which can get implemented on computers (time: 1970s)

A one-slide history of Neurons and Neuronal Networks

- 1940's McCulloch and Pitts Neuronal Description
Donald Hebb "The Organization of Behaviour"
- late 1960's Rosenblatt's Perceptron
- 1970's Minsky and Papert's criticism (1969), first *AI winter*.
Concerns and unsolved obstacles led to lower interest and poor funding
- 1980's Hopfield's impact (self-learning)
- 1990's training through error-backpropagation (1986)
- 2000's more mathematically rigorous statistical learning theory
- 2000's more computing power → deep neural networks

Applications today: Pattern recognition: classification and clustering; General interpolation problems; Data representations, coding and compression; Signal processing; Time series prediction; System identification; Decision support (e.g. medical or industrial diagnostics); Memory storage, modelling; Optimization, combinatorial problems

What can a Single Neuron do?



The **net input** of the neuron is given by


$$\text{net} = \sum_{i=1}^n p_i w_i = p_1 w_1 + p_2 w_2 + \cdots + p_n w_n$$

whereas the output of the neuron is computed as


$$\text{out} = f(\text{net} - \theta)$$

... anything ... depending on the activation function $f(\cdot)$


Typical Activation Functions $f(\cdot)$




Linear function




Step function




Ramp function



Sigmoid function



Hyperbolic tangent function



Gaussian function

By hand-selecting the activation function, we can allow arbitrary functionality.



But we can't know a-priori!

Hence, we require a generic system that works independent of activation function.

What can a Single Neuron do?

Assume a **LINEAR** activation function. “Boolean Logic” Classification on 2-dimensional input, decision based on $\text{out} \geq 0$ or $\text{out} < 0$

AND problem





Sample solution:

$$w_1 = w_2 = 0.8$$

$$\theta = 1$$

OR problem




Sample solution:

$$w_1 = w_2 = 0.8$$

$$\theta = 0.5$$

XOR problem



Sample solution:

The problem is not linearly separable →

we need a **multi-layer perceptron**



Training Neurons Networks

SUPERVISED LEARNING (labeled training data exists)

Idea: use the error produced by a neuron to adjust its weights

Error metric of a neuron

$$E(t) = \sum_{p=1}^{P_T} \left(t_p(t) - out_p(t) \right)^2$$

t_p target output
 $out_p(t)$ current output

Weight update for a neuron i

$$w_i(t) = w_i(t - 1) + \Delta w_i(t)$$


$$\Delta w_i(t) = \eta \left(\frac{-\partial E(t)}{\partial w_i(t)} \right)$$

This updates all neuronal input weights w_i for a all data point p , subject to the current weights.

Training a single neuron, real-world example

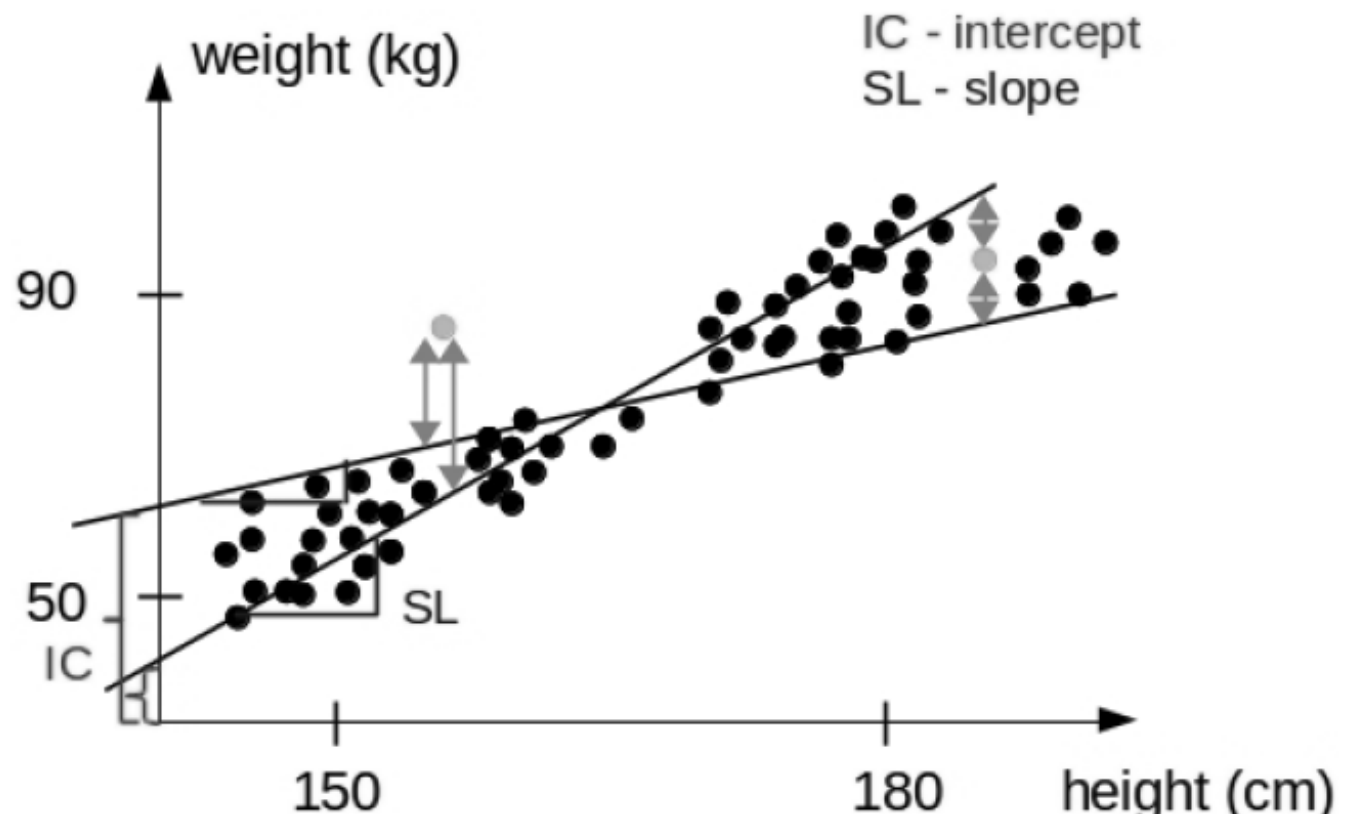
height

Bias offset
(= const 1)




$$\begin{aligned}\text{weight} &= w1 * \text{height} + w2 * \text{bias} \\ &= SL * \text{height} + IC\end{aligned}$$

Every new data point will adjust **SL** and **IC** such that ultimately the best matching linear approximation is encoded in w_1, w_2




Training a single neuron, real-world example




How do we achieve this? In order to guide the search for the suitable parameters Gradient Descent is used:

1. Pick **random initial values** for IC/SL (e.g. x_0)
2. **Calculate the gradient** with respect to each model parameter (i.e. IC, SL)
3. **Update the parameters** in the **direction** of the **negative gradient**
4. **Repeat 2 and 3 until convergence** (e.g. $x_1 \dots x_4$)




From Neurons to Networks

Single neuron "node"




Neuronal Networks


Single Layer






Multi Layer



Recurrent



Neuronal Networks




Main Challenges

- Which structure to use?
- How to train (adjust weights) given a particular problem?
We learned how to do for one neuron, but for networks?
- (How to provide sufficient computing power?)

How to train Neuronal Networks

- We can not simply train neurons in a network, as the “desired output” (t_p) is unknown for neurons in middle layers
- Instead, **ERROR BACKPROGATION** through the network



$$\Delta w_{ik}(t) = f'(net_i(t)) \cdot \sum_{j \in P} \delta_j(t) \cdot w_{ji}(t) \cdot f(net_k(t))$$

$$\delta_i(t) = f'(net_i(t)) \cdot \sum_{j \in P} \delta_j(t) \cdot w_{ji}(t)$$

Tricks of the Trade – how to get your network working

1. Weight Initialization

initialize all weights to **small random** numbers

2. Input / Output Normalization

output: use linear output neuron

input: try to normalize to $[-1 \dots +1]$ or close to that

3. Use Small Weight Update η

Constant or time dependent $\eta(t)$.

4. Avoid local minima in weight space

Simulated Annealing, multiple independent runs

Tricks of the Trade – how to get your network working


5. Network size

Complete mystery; rules of thumb, fan-in / fan-out (projections)


6. Vanishing Gradient

multiple layers significantly reduce the available gradient

7. Overtraining / Overfitting




Overtraining / Overfitting



Split available data in two sets

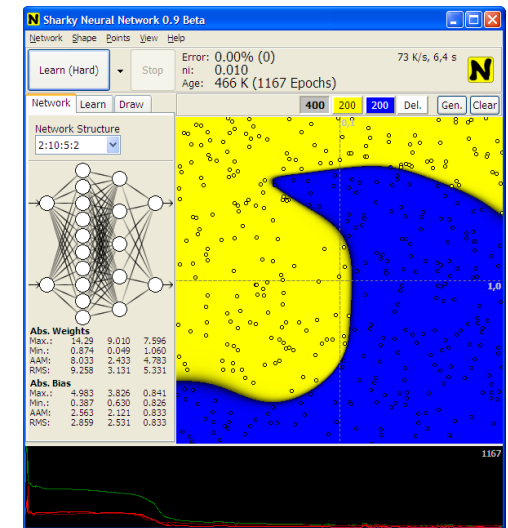
-) 70% Training Data
-) 30% Validation

Train on training data only,
evaluate on Validation data,
STOP training as performance drops



Software Example Neuronal Nets

- Google for “Neuronal Network Online Software Simulation” or similar
- http://www.sharktime.com/en_SharkyNeuralNetwork.html



- https://grey.colorado.edu/emergent/index.php/Comparison_of_Neural_Network_Simulators



Other Forms of Neuronal Networks

- Recurrent Neuronal Networks
 - Liquid State Machine
- Unsupervised Learning
 - Hebbian Networks
 - Kohonan Self-Organizing Maps
- Deep Neuronal Networks
 - Deeeeeeeep

more interest in this topic?

KTH DD2437 – Artificial Neural Networks and Deep Architectures

A Brief Introduction to Our Own Research



Computation in Brains

Getting to know your Brain

- 1.3 Kg, about 2% of body weight
- 10^{11} neurons
- Neuron growth
250.000 / min (early pregnancy)
... but also loss, 1 neuron/second

“Operating Mode” of Neurons

- Analog leaky integration in soma
- Digital pulses (spikes) along neurites
- 10^{14} stochastic synapses
- Typical operating “frequency”
 $\leq 100\text{Hz}$, typically $\sim 10\text{Hz}$, asynchronous

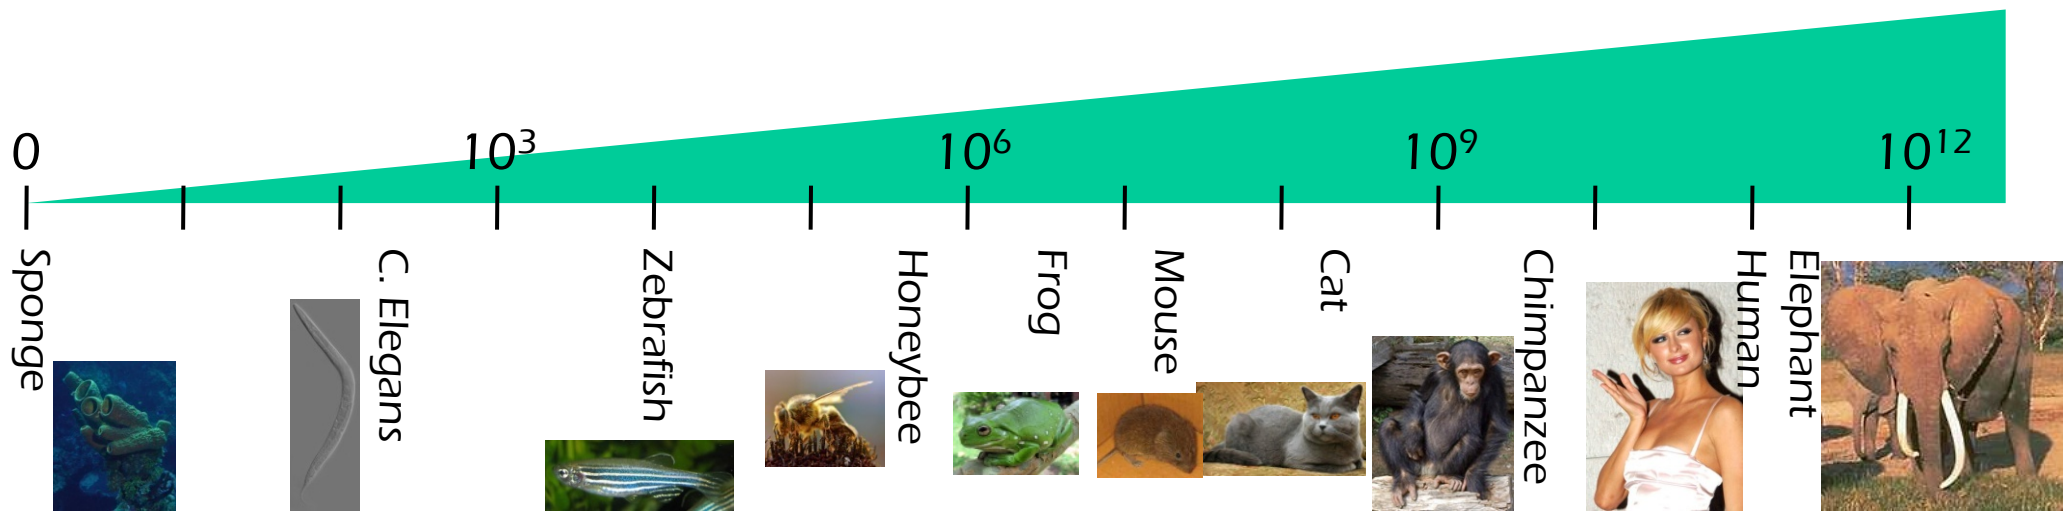
... and in Machines

Getting to know your Computer’s CPU

- 50g, largely irrelevant
- 10^{10} transistors (SPARC M7, 2015)
- ideally no modification over lifetime

“Operating Mode” of CPUs

- Digital Boolean logic processing
- Digital signal propagation
- Reliable remote storage of data
- Typical operating frequency
Several GHz, synchronous




Real-Time Neuronal Information Processing in Closed-Loop Systems

- Event-Based Neuromorphic Vision
- Information Processing in Distributed Neuronal Circuits
- Neuromorphic Real-Time Computing and Control
- Self-Construction and Organization of Neuronal Circuits




Event-Based
Neuromorphic Vision

Neuromorphic Real-Time
Computing and Control



Information Processing in
Distributed Neuronal Circuits




Self-Construction and
Organization of Neuronal Circuits

Neuro Computing Hardware Systems




CPU/GPU

- Software Neuronal Net Simulation
- High Energy
- Central Code and Data Memory
- Limited Communication Bandwidth
- Limited Parallelism
- High System Clock




SpiNNaker

- Software Neuronal Net Simulation
- Reduced Energy
- Distributed Code and Data Memory
- Optimized for Spiking Communication
- Massive Parallelism
- Reduced System Clock



TrueNorth

- Neuronal Circuits in Digital Hardware
- Low Energy
- Function given by Synaptic Connectivity
- Limited Local Spiking Communication
- Expandable Parallelism
- Low System Clock



Spikey

- Neuronal Circuits in Analog Hardware
- Minimal Energy
- Function given by Synaptic Connectivity
- (Limited Local Spiking Communication)
- (Limited Parallelism)
- Real-time Accelerated System Clock